

Project Part C: Classification



```
In [57]: analyst = "Khoa Nguyen" # Replace this with your name
```

```
In [58]: f = "setup.R"; for (i in 1:10) { if (file.exists(f)) break else f = paste0("../", f) }; source(f)
options(repr.matrix.max.rows=674)
```

1 Introduction

1.1 Objective

Build, evaluate, and tune a classifier trained on a transformed dataset about public company fundamentals. Later, use the classifier along with additional analysis to recommend a portfolio of 12 company investments that maximizes 12-month return of an overall

1.2 Approach

Retrieve a dataset ready for predictive model construction.

Build a model to predict whether stock price will grow more than 30% over 12 months, given 12 months of past company fundamentals data, using a machine learning model construction method.

Evaluate and tune the model for optimal business performance.

2 Business Model & Business Parameters

The business model is ...

$$\text{profit} = \left(\sum_{i \in \text{portfolio}} (1 + \text{growth}_i) \times \text{allocation}_i \right) - \text{budget}$$

$$\text{profit rate} = \text{profit} \div \text{budget}$$

$$\text{budget} = \sum_{i \in \text{portfolio}} \text{allocation}_i$$

Business parameters include ...

- budget is total investment to allocate across the companies in the portfolio
- portfolio size is number of companies in the portfolio
- allocation is vector of amounts to allocate to specific companies in the portfolio, must sum to budget
- threshold is growth that qualifies as lowest attractive growth

In [59]: *# Set the business parameters.*

```
budget = 1000000
portfolio_size = 12
allocation = rep(budget/portfolio_size, portfolio_size)

fmtsx(fmt(budget), fmt(portfolio_size), fmt(allocation))
```

<u>budget</u>	<u>portfolio_size</u>	<u>allocation</u>
1,000,000	12	83,333
		83,333
		83,333
		83,333
		83,333
		83,333
		83,333
		83,333
		83,333
		83,333
		83,333
		83,333

Portfolio to be filled with companies predicted to have the highest probabilities of growing above 30%.

3 Data

```
In [60]: # Retrieve data.
# How many observations and variables?
# Present the first few observations.

data = read.csv("My Data.csv", header=TRUE, na.strings=c("NA", ""), stringsAsFactors=FALSE)
data$big_growth = factor(data$big_growth, levels=c("YES", "NO"))

fmtx(size(data))
fmtx(data[1:6,], FFO)
```

size(data)

observations	variables
--------------	-----------

4,305	9
-------	---

data (first few observations)

big_growth	growth	prccq	gvkey	tic	conm	PC1	PC2	PC3
NO	0.0507	43.69	1,004	AIR	AAR CORP	1.4098	0.2125	-0.1874
NO	-0.3829	32.11	1,045	AAL	AMERICAN AIRLINES GROUP INC	-2.8093	0.2246	1.4366
YES	0.3158	6.75	1,050	CECE	CECO ENVIRONMENTAL CORP	1.5247	0.4396	-0.1679
NO	-0.2165	8.66	1,062	ASA	ASA GOLD AND PRECIOUS METALS	1.5737	0.6384	0.0123
NO	-0.1185	15.25	1,072	AVX	AVX CORP	1.2813	0.4529	0.0929
NO	0.0002	85.20	1,075	PNW	PINNACLE WEST CAPITAL CORP	0.3698	-0.4861	-0.0128

4 Build Classification Model

```
In [61]: model = naiveBayes(big_growth ~ PC1 + PC2 + PC3, data)
model
```

Naive Bayes Classifier for Discrete Predictors

Call:

```
naiveBayes.default(x = X, y = Y, laplace = laplace)
```

A-priori probabilities:

Y

	YES	NO
	0.08362	0.91638

Conditional probabilities:

PC1

Y	[,1]	[,2]
YES	1.129	1.332
NO	-0.103	5.674

PC2

Y	[,1]	[,2]
YES	0.24124	0.898
NO	-0.02201	4.793

PC3

Y	[,1]	[,2]
YES	-0.014239	0.677
NO	0.001299	3.653

5 Evaluate Classification Model (5-fold cross-validation)

```
In [62]: set.seed(0)
         fold = createFolds(data$big_growth, k=5)
         str(fold)
```

List of 5

```
$ Fold1: int [1:861] 9 13 17 19 31 42 44 54 60 66 ...
$ Fold2: int [1:861] 1 2 6 11 16 25 32 49 55 59 ...
$ Fold3: int [1:861] 4 8 14 22 28 34 40 45 50 52 ...
$ Fold4: int [1:861] 3 5 15 18 21 24 26 27 30 36 ...
$ Fold5: int [1:861] 7 10 12 20 23 29 33 35 37 46 ...
```

```

In [63]: # accuracy = c()
fold_performance = data.frame()

for (i in 1:5)
{ data.test = data[fold[[i]],]
  data.train = data[setdiff(1:nrow(data), fold[[i]]),]
  model_train = naiveBayes(big_growth ~ PC1 + PC2 + PC3, data.train)
  prob = predict(model_train, data.test, type="raw") # use predict arguments appropriate for your model
  class.predicted = as.class(prob, class="YES", cutoff=0.5)
  CM = confusionMatrix(class.predicted, data.test$big_growth)$table
  cm = CM/sum(CM)
  accuracy = cm[1,1]+cm[2,2]

  data.test$class.predicted = class.predicted
  data.test$prob = prob[,1]
  data.test = data.test[data.test$class.predicted == "YES",]
  data.sorted = data.test[order(data.test$prob, decreasing=TRUE),]
  company.data.growth = data.sorted[1:12, "growth"]
  profit = sum((1 + company.data.growth)*allocation) - budget
  fold_performance = rbind(fold_performance, data.frame(fold=i, accuracy=accuracy, profit = profit))}

fmtx(fold_performance, "Fold Performance")

```

Fold Performance

fold	accuracy	profit
1	0.2323	-144,476
2	0.2230	-114,764
3	0.2334	-22,672
4	0.2033	4,896
5	0.2021	-119,455

```
In [64]: accuracy.cv = mean(fold_performance$accuracy)
profit.cv = mean(fold_performance$profit)
profit_rate.cv = profit.cv/budget
fmtx(data.frame(accuracy.cv, profit.cv, profit_rate.cv), "5-Fold Cross-Validation Estimated Performance")
```

5-Fold Cross-Validation Estimated Performance

accuracy.cv	profit.cv	profit_rate.cv
0.2188	-79,294	-0.0793

6 Tune Classification Model

```
In [65]: tune = data.frame()
for (f in exhaustive(names(data[,c("PC1", "PC2", "PC3")]), keep="big_growth")) # try every combination of variables
for (q in c(0.25, 0.33, 0.50)) # try several values for cutoff
{
  nfold = 5
  set.seed(0)
  fold = createFolds(data$big_growth, k=nfold)
  accuracy = c()
  profit = c()
  for (i in 1:nfold) { data.train = data[setdiff(1:nrow(data), fold[[i]]),]
    data.test = data[fold[[i]],]
    model = naiveBayes(big_growth ~ ., data.train[,f], laplace=TRUE)
    prob = predict(model, data.test, type="raw")
    class.predicted = as.class(prob, class="YES", cutoff=q)
    CM = confusionMatrix(class.predicted, data.test$big_growth)$table
    cm = CM/sum(CM)
    accuracy[i] = cm[1,1]+cm[2,2]

    data.test$class.predicted = class.predicted
    data.test$prob = prob[,1]
    data.sorted = data.test[order(data.test$prob, decreasing=TRUE),]
    company.data.growth = data.sorted[1:12, "growth"]
    profit[i] = sum((1 + company.data.growth)*allocation) - budget }

  accuracy.cv = mean(accuracy)
  profit.cv = mean(profit)
  profit_rate.cv = profit.cv/budget

  tune = rbind(tune, data.frame(method="naive bayes", variables=vector2string(f), cutoff = q,
                                accuracy.cv, profit.cv, profit_rate.cv))
}
best = tune[which.max(tune$profit.cv),]
fmtx(best, "best model")

fmtx(tune, "search for best model")
```

best model					
method	variables	cutoff	accuracy.cv	profit.cv	profit_rate.cv
naive bayes	PC3, big_growth	0.25	0.2049	47,492	0.0475

search for best model

method	variables	cutoff	accuracy.cv	profit.cv	profit_rate.cv
naive bayes	PC1, big_growth	0.25	0.2987	-85,297	-0.0853
naive bayes	PC1, big_growth	0.33	0.8035	-85,297	-0.0853
naive bayes	PC1, big_growth	0.50	0.9134	-85,297	-0.0853
naive bayes	PC2, big_growth	0.25	0.3554	-146,897	-0.1469
naive bayes	PC2, big_growth	0.33	0.7022	-146,897	-0.1469
naive bayes	PC2, big_growth	0.50	0.9157	-146,897	-0.1469
naive bayes	PC3, big_growth	0.25	0.2049	47,492	0.0475
naive bayes	PC3, big_growth	0.33	0.7823	47,492	0.0475
naive bayes	PC3, big_growth	0.50	0.9122	47,492	0.0475
naive bayes	PC1, PC2, big_growth	0.25	0.2197	-142,451	-0.1425
naive bayes	PC1, PC2, big_growth	0.33	0.2355	-142,451	-0.1425
naive bayes	PC1, PC2, big_growth	0.50	0.3954	-142,451	-0.1425
naive bayes	PC1, PC3, big_growth	0.25	0.2007	-116,848	-0.1168
naive bayes	PC1, PC3, big_growth	0.33	0.2107	-116,848	-0.1168
naive bayes	PC1, PC3, big_growth	0.50	0.2381	-116,848	-0.1168
naive bayes	PC2, PC3, big_growth	0.25	0.1823	-112,989	-0.1130
naive bayes	PC2, PC3, big_growth	0.33	0.1954	-112,989	-0.1130
naive bayes	PC2, PC3, big_growth	0.50	0.2383	-112,989	-0.1130
naive bayes	PC1, PC2, PC3, big_growth	0.25	0.1991	-79,294	-0.0793
naive bayes	PC1, PC2, PC3, big_growth	0.33	0.2046	-79,294	-0.0793
naive bayes	PC1, PC2, PC3, big_growth	0.50	0.2188	-79,294	-0.0793

